

Neural Networks

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبدالله
للتكنولوجيا
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

November 26, 2025


Table of Content

- ▶ What a neural network is
- ▶ How it fits with loss and optimizer
- ▶ Forward and backward passes
- ▶ Activation functions
- ▶ Inputs and outputs for common tasks
- ▶ Optimizers used in deep learning
- ▶ Architectures

- ▶ A neural network is a function that maps input features to outputs using layers of simple units called neurons
- ▶ Each neuron computes a weighted sum of its inputs then applies a non linear activation
- ▶ Stacking layers lets the model learn complex patterns that linear models cannot capture

- ▶ **Input layer** holds the features
- ▶ **Hidden layers** learn intermediate representations
- ▶ **Output layer** produces predictions

Neural Network



► Neural network model

- Maps input x to prediction $\hat{y} = f_{\theta}(x)$
- θ are the parameters (weights and biases)


► Loss function

- Measures how far predictions are from targets
- Guides what it means to be a good model for the task

► Optimizer

- Uses gradients of the loss to update θ
- Tries to find parameters that reduce the loss on data

The Training Loop



Forward Pass:

- ▶ Data → Model → Predictions
- ▶ Compute Loss

Backward Pass:

- ▶ Calculate gradients
- ▶ Update parameters

Repeat until loss converges (model learns)!

The Forward Pass

- ▶ The forward pass starts from inputs and flows through the network
- ▶ Each layer applies a linear transformation then an activation
- ▶ At the end we compute the loss between predictions and targets

Example for one hidden layer

$$a^{(1)} = \sigma(W^{(1)}x + b^{(1)})$$

$$\hat{y} = g(W^{(2)}a^{(1)} + b^{(2)})$$

- ▶ $x \in \mathbb{R}^k$ input features
- ▶ $a^{(1)}$ hidden layer activations
- ▶ g is the output activation chosen for the task

The Backward Pass


- ▶ The backward pass starts from the loss and flows backward through the network
- ▶ It uses the chain rule to compute gradients of the loss with respect to every parameter
- ▶ These gradients tell the optimizer how to change weights to reduce the loss

At a high level


- ▶ Forward pass gives \hat{y} and loss $J(\theta)$
- ▶ Backward pass gives gradients $\nabla_{\theta}J$
- ▶ Optimizer updates parameters using these gradients

Forward and Backward Together

Forwardpass



Backwardpass



Single Neuron Computation

- ▶ Input features $x \in \mathbb{R}^k$
- ▶ Weights $w \in \mathbb{R}^k$ and bias $b \in \mathbb{R}$

$$z = w^\top x + b \quad , \quad a = \sigma(z)$$

- ▶ z is the pre activation (linear part)
- ▶ σ is an activation function
- ▶ a is the neuron output that feeds the next layer or the loss

Why We Need Activation Functions

- ▶ If we only stack linear layers without activations we still get a linear function
- ▶ Non linear activations let the network approximate complex functions
- ▶ They introduce bends and thresholds that help separate classes and model non linear trends

Common Activation Functions

ReLU

$$\sigma(z) = \max(0, z)$$

- ▶ Simple and works well in deep nets
- ▶ Keeps positive values and drops negatives

Sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- ▶ Outputs values in $(0, 1)$
- ▶ Used for probabilities in binary classification

Tanh

$$\sigma(z) = \tanh(z)$$

- ▶ Outputs values in $(-1, 1)$
- ▶ Often used in recurrent networks

Softmax for Multiclass Outputs

- ▶ For multiclass classification the output layer gives scores $z \in \mathbb{R}^C$ for C classes
- ▶ The softmax function turns scores into a probability distribution

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

- ▶ Each output is between 0 and 1
- ▶ All outputs sum to 1

Inputs and Outputs in Practice

- ▶ Input vector $x \in \mathbb{R}^k$
 - Example features
 - Age income number of past purchases
- ▶ Output shape depends on the task
 - Scalar for single output regression or binary classification
 - Vector for multi output regression or multiclass classification

Network Shapes for Different Tasks

Task	Input	Output of NN	Typical loss
Single output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}$	Mean squared error
Multi output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}^m$	MSE over all outputs
Binary classification	$x \in \mathbb{R}^k$	$\hat{p} \in (0, 1)$	Binary cross entropy
Multiclass classification	$x \in \mathbb{R}^k$	$\hat{p} \in \mathbb{R}^C$ softmax	Cross entropy

Examples of Heads on Top of the Same Body

- ▶ The hidden layers can be shared across tasks
- ▶ Only the last layer and the loss change
- ▶ Regression head
 - Last layer has one neuron with linear activation
 - Use mean squared error
- ▶ Binary head
 - Last layer has one neuron with sigmoid
 - Use binary cross entropy
- ▶ Multiclass head
 - Last layer has C neurons with softmax
 - Use cross entropy over classes

Optimizers: Gradient Descent

- ▶ Parameters θ are updated in the direction that reduces the loss

$$\theta_{\text{new}} = \theta_{\text{old}} - \alpha \nabla_{\theta} J(\theta)$$

- ▶ α is the learning rate
- ▶ $\nabla_{\theta} J(\theta)$ is the gradient of the loss with respect to the parameters
- ▶ In deep learning we usually use mini batch gradient descent

AdaGrad

- ▶ Keeps a running sum of squared gradients for each parameter:

$$G_t = G_{t-1} + g_t^2$$

- ▶ Update rule with parameter wise learning rate:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t} + \epsilon} g_t$$

- ▶ Parameters with large accumulated gradients get smaller effective learning rates
- ▶ Works well for sparse features

RMSProp

- ▶ Uses an exponential moving average of squared gradients:

$$E[g^2]_t = \alpha E[g^2]_{t-1} + (1 - \alpha) g_t^2$$

- ▶ Update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

- ▶ Prevents the learning rate from shrinking too fast and keeps training going
- ▶ Often used as a base idea in modern optimizers

Adam

- ▶ Keeps a moving average of gradients and squared gradients:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

- ▶ Bias correction:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- ▶ Update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

- ▶ Combines momentum + adaptive step size and usually speeds up training

AdamW


- ▶ Uses the same Adam update, plus **decoupled weight decay**:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t - \eta \lambda \theta_t$$


- ▶ Weight decay is applied directly to parameters instead of through the gradient
- ▶ Often used as a strong default optimizer for modern deep networks

Types of Neural Network Architectures


- ▶ **MLP** fully connected network for tabular data
- ▶ **CNN** convolutional neural network for images
- ▶ **RNN and variants** for sequences and time series



Standard NN



Convolutional NN




Recurrent NN

What Is an Autoencoder

- ▶ An autoencoder is a network that tries to reconstruct its input
- ▶ It has two parts
 - **Encoder** maps x to a lower dimensional code z
 - **Decoder** maps z back to a reconstruction \hat{x}
- ▶ The loss compares \hat{x} and x usually with mean squared error

Autoencoder Diagram



The bottleneck forces the network to learn a compressed representation of the data

Autoencoder Applications

- ▶ Dimensionality reduction and visualization
- ▶ Denoising images or signals
- ▶ Anomaly detection using high reconstruction error
- ▶ As a pretraining step to learn useful representations

- ▶ Andrew Ng, Deep Learning Specialization and Machine Learning course
- ▶ Aurélien Géron, Hands On Machine Learning with Scikit Learn Keras and TensorFlow
- ▶ Ian Goodfellow Yoshua Bengio Aaron Courville, Deep Learning
- ▶ DeepLearning.AI short courses on optimization and advanced architectures