

# Neural Networks

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبدالله  
للتكنولوجيا  
King Abdullah University of  
Science and Technology

KAUST Academy  
King Abdullah University of Science and Technology


November 29, 2025

# Table of Content

- ▶ What a neural network is
- ▶ How it fits with loss and optimizer
- ▶ Forward and backward passes
- ▶ Activation functions
- ▶ Inputs and outputs for common tasks
- ▶ Optimizers used in deep learning
- ▶ Architectures


- ▶ A neural network is a function that maps input features to outputs using layers of simple units called neurons
- ▶ Each neuron computes a weighted sum of its inputs then applies a non linear activation
- ▶ Stacking layers lets the model learn complex patterns that linear models cannot capture
  
- ▶ **Input layer** holds the features
- ▶ **Hidden layers** learn intermediate representations
- ▶ **Output layer** produces predictions

# Neural Network



## Why Multiple Layers? (Depth)

- ▶ A single hidden layer network is already very powerful in theory
  - ▶ In practice, **deeper** networks:
    - Construct more powerful intermediate features
    - Can represent complex functions more **efficiently** (fewer neurons overall)



## ► Neural network model

- Maps input  $x$  to prediction  $\hat{y} = f_{\theta}(x)$
- $\theta$  are the parameters (weights and biases)


## ► Loss function

- Measures how far predictions are from targets
- Guides what it means to be a good model for the task

## ► Optimizer

- Uses gradients of the loss to update  $\theta$
- Tries to find parameters that reduce the loss on data

# The Training Loop



## Forward Pass:

- ▶ Data → Model → Predictions
- ▶ Compute Loss

## Backward Pass:

- ▶ Calculate gradients
- ▶ Update parameters

*Repeat until loss converges (model learns)!*

# The Forward Pass

- ▶ The forward pass starts from inputs and flows through the network
- ▶ Each layer applies a linear transformation then an activation
- ▶ At the end we compute the loss between predictions and targets

## Example for one hidden layer

$$a^{(1)} = \sigma(W^{(1)}x + b^{(1)})$$

$$\hat{y} = g(W^{(2)}a^{(1)} + b^{(2)})$$

- ▶  $x \in \mathbb{R}^k$  input features
- ▶  $a^{(1)}$  hidden layer activations
- ▶  $g$  is the output activation chosen for the task

# The Backward Pass


- ▶ The backward pass starts from the loss and flows backward through the network
- ▶ It uses the chain rule to compute gradients of the loss with respect to every parameter
- ▶ These gradients tell the optimizer how to change weights to reduce the loss

## At a high level


- ▶ Forward pass gives  $\hat{y}$  and loss  $J(\theta)$
- ▶ Backward pass gives gradients  $\nabla_{\theta}J$
- ▶ Optimizer updates parameters using these gradients

# Forward and Backward Together

Forwardpass



Backwardpass



# Single Neuron Computation

- ▶ Input features  $x \in \mathbb{R}^k$
- ▶ Weights  $w \in \mathbb{R}^k$  and bias  $b \in \mathbb{R}$

$$z = w^\top x + b \quad , \quad a = \sigma(z)$$

- ▶  $z$  is the pre activation (linear part)
- ▶  $\sigma$  is an activation function
- ▶  $a$  is the neuron output that feeds the next layer or the loss

# Why We Need Activation Functions

- ▶ If we only stack linear layers without activations we still get a linear function
- ▶ Non linear activations let the network approximate complex functions
- ▶ They introduce bends and thresholds that help separate classes and model non linear trends

# Common Activation Functions

## ReLU

$$\sigma(z) = \max(0, z)$$

- ▶ Simple and works well in deep nets
- ▶ Keeps positive values and drops negatives

## Sigmoid

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- ▶ Outputs values in  $(0, 1)$
- ▶ Used for probabilities in binary classification

## Tanh

$$\sigma(z) = \tanh(z)$$

- ▶ Outputs values in  $(-1, 1)$
- ▶ Often used in recurrent networks

# Softmax for Multiclass Outputs

- ▶ For multiclass classification the output layer gives scores  $z \in \mathbb{R}^C$  for  $C$  classes
- ▶ The softmax function turns scores into a probability distribution

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

- ▶ Each output is between 0 and 1
- ▶ All outputs sum to 1

# Inputs and Outputs in Practice

- ▶ Input vector  $x \in \mathbb{R}^k$ 
  - Example features
  - Age income number of past purchases
- ▶ Output shape depends on the task
  - Scalar for single output regression or binary classification
  - Vector for multi output regression or multiclass classification

# Network Shapes for Different Tasks

Task	Input	Output of NN	Typical loss
Single output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}$	Mean squared error
Multi output regression	$x \in \mathbb{R}^k$	$\hat{y} \in \mathbb{R}^m$	MSE over all outputs
Binary classification	$x \in \mathbb{R}^k$	$\hat{p} \in (0, 1)$	Binary cross entropy
Multiclass classification	$x \in \mathbb{R}^k$	$\hat{p} \in \mathbb{R}^C$ softmax	Cross entropy

# Examples of Heads on Top of the Same Body

- ▶ The hidden layers can be shared across tasks
- ▶ Only the last layer and the loss change
- ▶ Regression head
  - Last layer has one neuron with linear activation
  - Use mean squared error
- ▶ Binary head
  - Last layer has one neuron with sigmoid
  - Use binary cross entropy
- ▶ Multiclass head
  - Last layer has  $C$  neurons with softmax
  - Use cross entropy over classes

# Optimizers: Gradient Descent

- ▶ Parameters  $\theta$  are updated in the direction that reduces the loss

$$\theta_{\text{new}} = \theta_{\text{old}} - \alpha \nabla_{\theta} J(\theta)$$

- ▶  $\alpha$  is the learning rate
- ▶  $\nabla_{\theta} J(\theta)$  is the gradient of the loss with respect to the parameters
- ▶ In deep learning we usually use mini batch gradient descent

## AdaGrad

- ▶ Keeps a running sum of squared gradients for each parameter:

$$G_t = G_{t-1} + g_t^2$$

- ▶ Update rule with parameter wise learning rate:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{G_t} + \epsilon} g_t$$

- ▶ Parameters with large accumulated gradients get smaller effective learning rates
- ▶ Works well for sparse features

## RMSProp

- ▶ Uses an exponential moving average of squared gradients:

$$E[g^2]_t = \alpha E[g^2]_{t-1} + (1 - \alpha) g_t^2$$

- ▶ Update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

- ▶ Prevents the learning rate from shrinking too fast and keeps training going
- ▶ Often used as a base idea in modern optimizers

## Adam

- ▶ Keeps a moving average of gradients and squared gradients:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

- ▶ Bias correction:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- ▶ Update rule:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

- ▶ Combines momentum + adaptive step size and usually speeds up training

## AdamW

- ▶ Uses the same Adam update, plus **decoupled weight decay**:

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t - \eta \lambda \theta_t$$

- ▶ Weight decay is applied directly to parameters instead of through the gradient
- ▶ Often used as a strong default optimizer for modern deep networks

## ▶ Learning rate $\eta$

- Too small: training is very slow
- Too large: loss may bounce or diverge

## ▶ Batch size

- Number of examples per parameter update
- Small batch: noisy but often good generalization
- Large batch: smoother gradients, needs more memory

## ▶ Epochs

- One pass over the whole training dataset (consists of batches/steps)

## ► Model side

- Number of layers and neurons per layer
- Activation functions (ReLU, sigmoid, tanh, ...)
- Output layer shape and activation

## ► Training side

- Loss function (MSE, cross entropy, ...)
- Optimizer (SGD, Adam, AdamW, ...)
- Hyperparameters (learning rate, batch size, epochs)

► **Underfitting:**

- Model is too simple or not trained enough

► **Overfitting:**

- Model memorizes training data

# How Can We Detect Overfitting?

- ▶ Idea: measure performance on **data the model did not see during training**
- ▶ We split our dataset:
  - A part for **training** the model
  - A part for **validation/testing** to check generalization
- ▶ If the model keeps improving on the training set but stops improving (or gets worse) on the validation set, it is overfitting.


# Hold Out Split: Train, Validation, Test

- ▶ We split the dataset into **disjoint** parts
  - **Train set** to fit the model
  - **Validation set** to tune models and hyperparameters
  - **Test set** kept aside for a final unbiased estimate
- ▶ For simplicity, people usually just split into train and test.
- ▶ We usually **shuffle** the data rows before splitting
  - So that all splits follow the same distribution as the full data
  - To avoid easy leakage such as first half of customers in train and last half in test

# Training vs Validation Loss


## Underfit

Model too simple




## Good Fit

Optimal balance



## Overfit

Memorizing data



**Train MSE:** 2.2  
**Val MSE:** 2.3

**Train MSE:** 0.6  
**Val MSE:** 0.7

**Train MSE:** 0.3  
**Val MSE:** 1.4



Train loss



Validation loss

# Reading Loss Curves

- ▶ During training we usually plot:
  - Training loss vs epochs
  - Validation loss vs epochs

*These curves are one of the most useful debugging tools in deep learning.*

# How To Fix Underfitting and Overfitting

## If the model is underfitting

- ▶ Use a more complex model
- ▶ Train longer
- ▶ Increase the number of layers and neurons

## If the model is overfitting

- ▶ Simplify the model
- ▶ Get more data
- ▶ Add regularization

- ▶ **Goal:** Keep the model powerful but prevent overfitting

## L2 Regularization (Weight Decay)

- ▶ Add a penalty on large weights to the loss:

$$J_{\text{total}}(\theta) = J(\theta) + \lambda \|\theta\|_2^2$$


- ▶ Encourages **smaller, smoother weights** that generalize better

## Early Stopping


- ▶ Monitor validation loss during training
- ▶ Stop training when validation loss stops improving
- ▶ Simple and very effective in practice

# Regularization: Dropout

- ▶ During training, **randomly turn off** (drop) some neuron outputs
- ▶ Each mini-batch sees a slightly different subnet of the network
- ▶ Intuition: the network cannot rely on a single neuron and learns more **robust** features
- ▶ At test time, we use the **full** network but scale activations appropriately




(a) Standard Neural Net




(b) After applying dropout.

# Types of Neural Network Architectures


- ▶ **MLP** fully connected network for tabular data
- ▶ **CNN** convolutional neural network for images
- ▶ **RNN and variants** for sequences and time series



Standard NN



Convolutional NN




Recurrent NN

# What Is an Autoencoder

- ▶ An autoencoder is a network that tries to reconstruct its input
- ▶ It has two parts
  - **Encoder** maps  $x$  to a lower dimensional code  $z$
  - **Decoder** maps  $z$  back to a reconstruction  $\hat{x}$
- ▶ The loss compares  $\hat{x}$  and  $x$  usually with mean squared error

# Autoencoder Diagram



The bottleneck forces the network to learn a compressed representation of the data

# Autoencoder Applications

- ▶ Dimensionality reduction and visualization
- ▶ Denoising images or signals
- ▶ Anomaly detection using high reconstruction error
- ▶ As a pretraining step to learn useful representations

# Autoencoder Example: Anomaly Detection

- ▶ Train an autoencoder to reconstruct **normal** examples (e.g., normal transactions)
- ▶ At test time:
  - Feed an input  $x$  through the autoencoder to get  $\hat{x}$
  - Compute reconstruction error:  $\|x - \hat{x}\|$
- ▶ If error is **small**: input looks similar to training data (likely normal)
- ▶ If error is **large**: input is unusual (possible anomaly)
- ▶ Very useful when we have many normal samples and few labeled anomalies

- ▶ Andrew Ng, Deep Learning Specialization and Machine Learning course
- ▶ Aurélien Géron, Hands On Machine Learning with Scikit Learn Keras and TensorFlow
- ▶ Ian Goodfellow Yoshua Bengio Aaron Courville, Deep Learning
- ▶ DeepLearning.AI short courses on optimization and advanced architectures