

# Recurrent Neural Networks (RNNs) Improvements

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبد الله  
للعلوم والتكنولوجيا  
King Abdullah University of  
Science and Technology

KAUST Academy

King Abdullah University of Science and Technology

# Table of Contents

1. Motivation
2. Learning Outcomes
3. Vanishing and Exploding Gradients
4. LSTM – Long Short-Term Memory
5. GRU – Gated Recurrent Unit
6. LSTM vs GRU
7. Use Cases for GRU/LSTM
8. Limitations of GRU/LSTM
9. Future Directions
10. Summary
11. References

## Why Do We Need to Improve RNNs?

- ▶ Vanilla RNNs **fail at capturing long-term dependencies**
- ▶ Gradients either **vanish or explode** over long sequences
- ▶ This **limits learning** over time-based tasks like translation, conversation modeling, or video understanding

**Solution:** Modify RNN architecture to retain important past information without instability.

# Learning Outcomes

By the end of this session, you should be able to:

- ▶ Explain why RNNs suffer from vanishing and exploding gradients
- ▶ Understand how LSTMs and GRUs solve these problems
- ▶ Compare LSTMs and GRUs in terms of performance and complexity
- ▶ Identify practical scenarios where each is preferred
- ▶ Recognize limitations and future directions in sequence modeling

# Vanishing and Exploding Gradients

**Backpropagation Through Time (BPTT)** spreads gradients across many time steps.

## Vanishing Gradients:

$$\left\| \frac{\partial L}{\partial h_t} \right\| \rightarrow 0$$

- ▶ Early layers barely learn
- ▶ Forget long-term dependencies

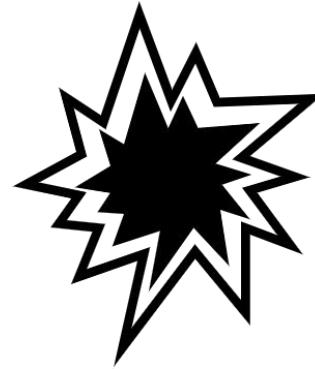
## Exploding Gradients:

$$\left\| \frac{\partial L}{\partial h_t} \right\| \rightarrow \infty$$

- ▶ Unstable updates, diverging weights

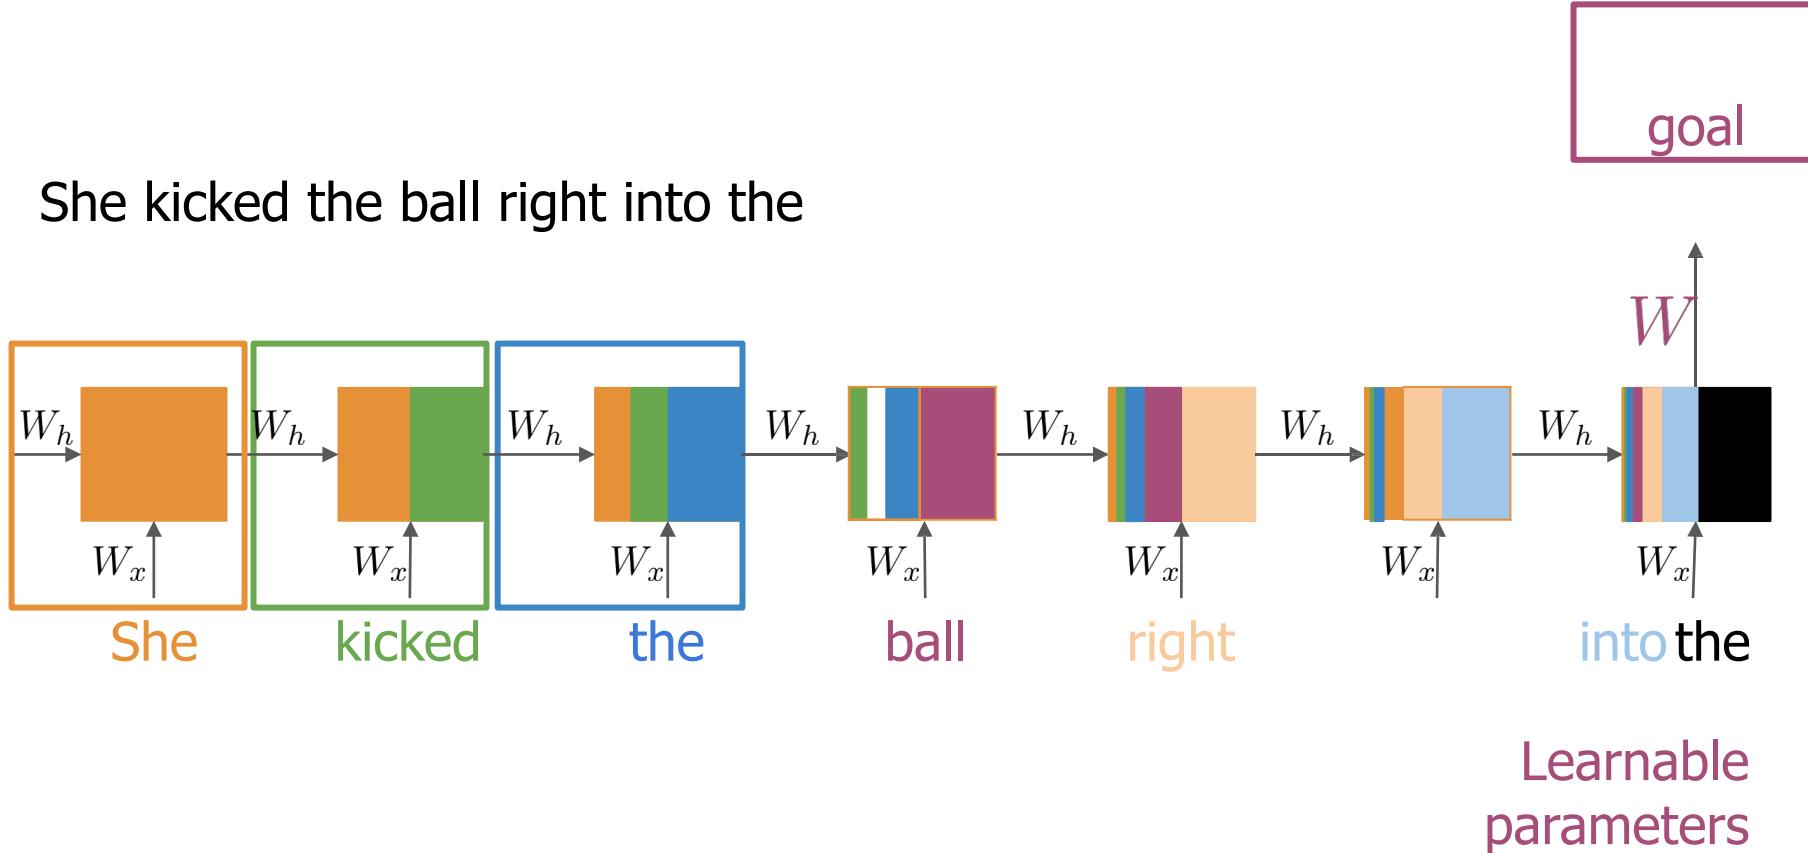
# RNNs and Vanishing Gradients

- Backprop through time
- RNNs and vanishing/exploding gradients
- Solutions

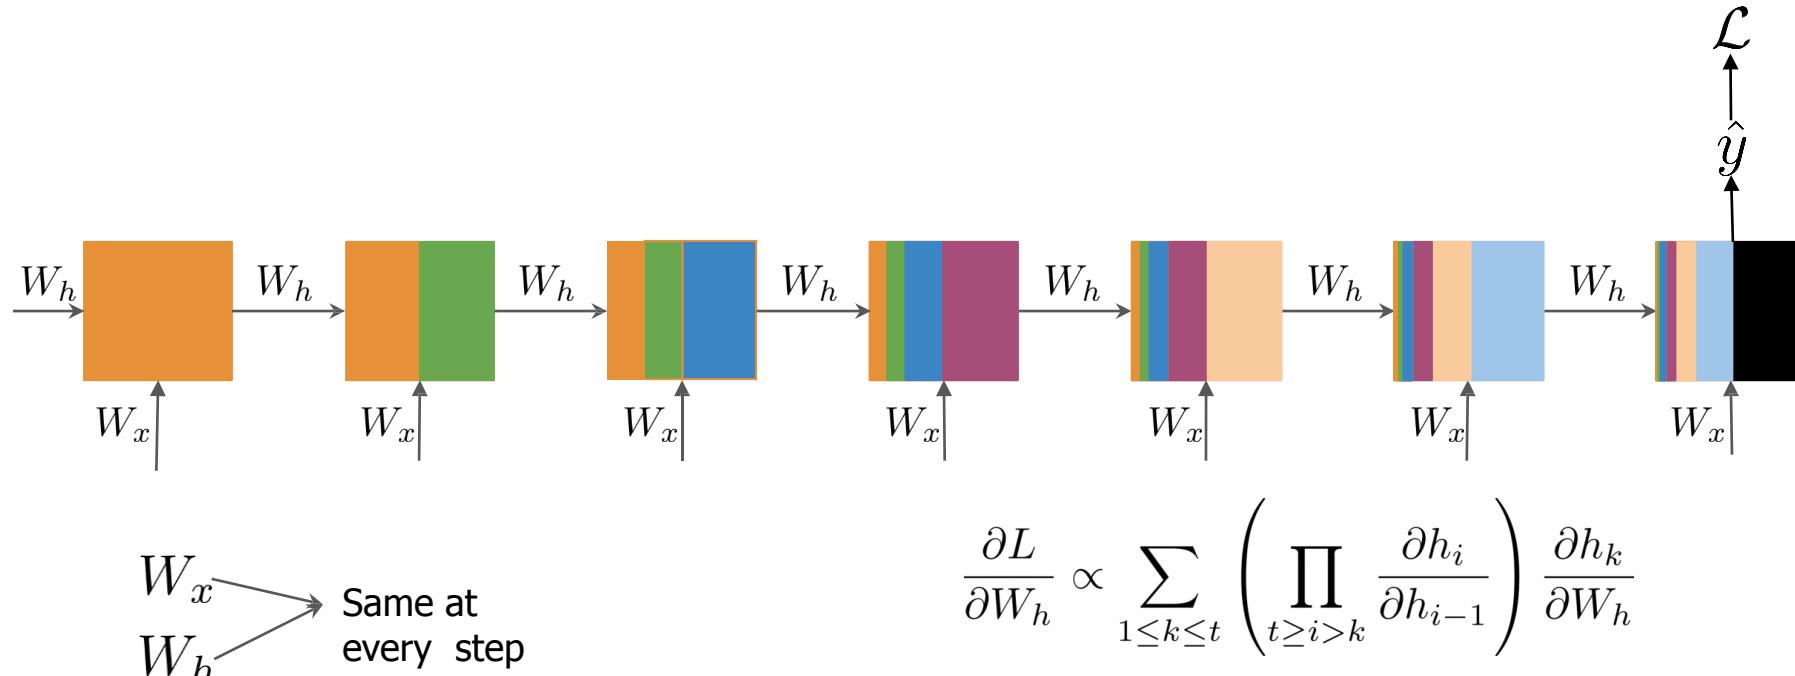


# RNN Basic Structure

She kicked the ball right into the



# Backpropagation through time



Gradient is proportional to a sum of  
partial derivative products

# Backpropagation through time

$$\frac{\partial L}{\partial W_h} \propto \sum_{1 \leq k \leq t} \left( \prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W_h}$$

Contribution of hidden state  $k$

Length of the product proportional to  
how far  $k$  is from  $t$

$$\frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial h_{t-2}} \frac{\partial h_{t-2}}{\partial h_{t-3}} \frac{\partial h_{t-3}}{\partial h_{t-4}} \frac{\partial h_{t-4}}{\partial h_{t-5}} \frac{\partial h_{t-5}}{\partial h_{t-6}} \frac{\partial h_{t-6}}{\partial h_{t-7}} \frac{\partial h_{t-7}}{\partial h_{t-8}} \frac{\partial h_{t-8}}{\partial h_{t-9}} \frac{\partial h_{t-9}}{\partial h_{t-10}} \frac{\partial h_{t-10}}{\partial W_h}$$

Contribution of hidden state  $t-10$

# Backpropagation through time

$$\frac{\partial L}{\partial W_h} \propto \sum_{1 \leq k \leq t} \left( \prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}} \right) \frac{\partial h_k}{\partial W_h}$$

Contribution of hidden state  $k$

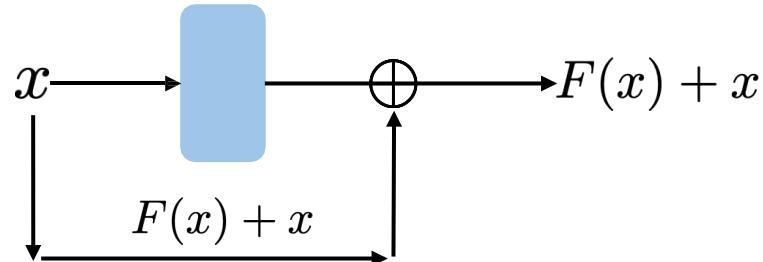
Length of the product proportional to  
how far  $k$  is from  $t$

Partial derivatives <1	Contribution goes to 0	Vanishing Gradient
Partial derivatives >1	Contribution goes to infinity	<b>Exploding</b> Gradient

# Solving for vanishing or exploding gradients

- Identity RNN with ReLU activation
- Gradient clipping
- Skip connections

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



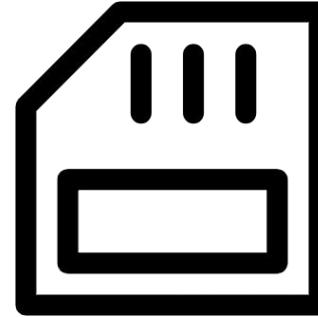
# RNNs: Advantages

- + Captures dependencies within a short range
- + Takes up less RAM than other n-gram models

# RNNs: Disadvantages

- Struggles to capture long term dependencies
- Prone to vanishing or exploding gradients

- Meet the Long short-term memory unit!
- LSTM architecture
- Applications



Introduced by Hochreiter & Schmidhuber (1997)

**Core Idea:** LSTM uses **gates** to control what to keep, forget, and output.

**Key Components:**

- ▶ **Forget Gate**  $f_t$
- ▶ **Input Gate**  $i_t$
- ▶ **Cell State**  $C_t$
- ▶ **Output Gate**  $o_t$

# LSTM – Long Short-Term Memory

## Equations:

$$f_t = \sigma(W_f[x_t, h_{t-1}] + b_f)$$

$$i_t = \sigma(W_i[x_t, h_{t-1}] + b_i)$$

$$\tilde{C}_t = \tanh(W_C[x_t, h_{t-1}] + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma(W_o[x_t, h_{t-1}] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

Helps retain long-term dependencies

# LSTMs: a memorable solution

- Learns when to remember and when to forget
- Basic anatomy:
  - A cell state
  - A hidden state
  - Multiple gates
- Gates allow gradients to avoid vanishing and exploding

# LSTMs: Based on previous understanding

Starting point with some irrelevant information



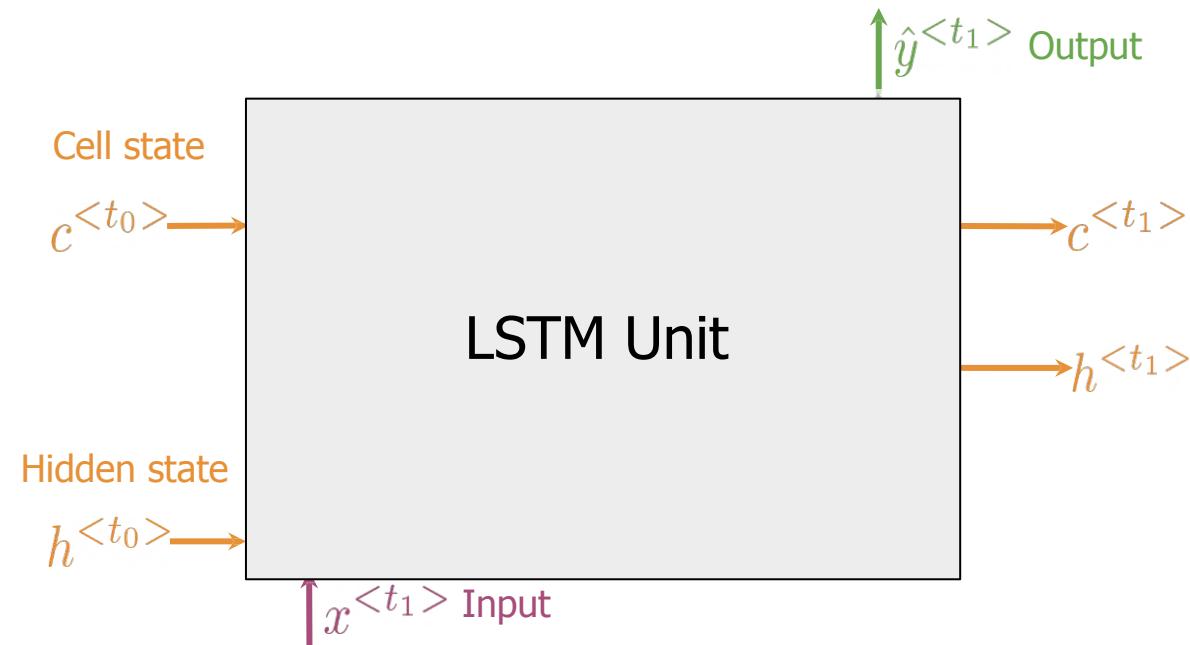
Cell and Hidden States

Discard anything irrelevant  
Add important new information  
Produce output

→ Gates



# Gates in LSTM

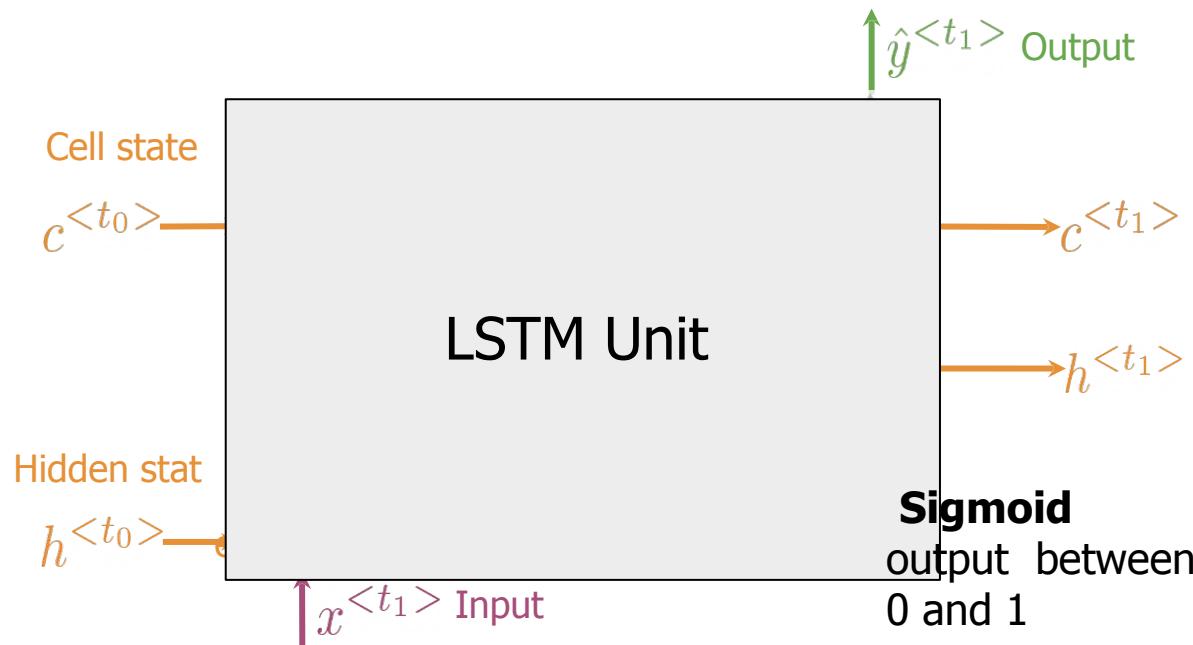


**1. Forget Gate:**  
information that is no longer important

**2. Input Gate:**  
information to be stored

**3. Output Gate:**  
information to use at current step

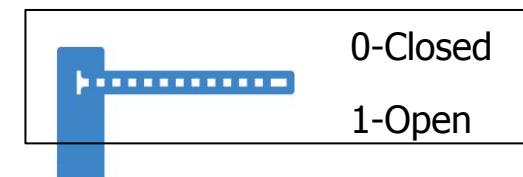
# Gates in LSTM



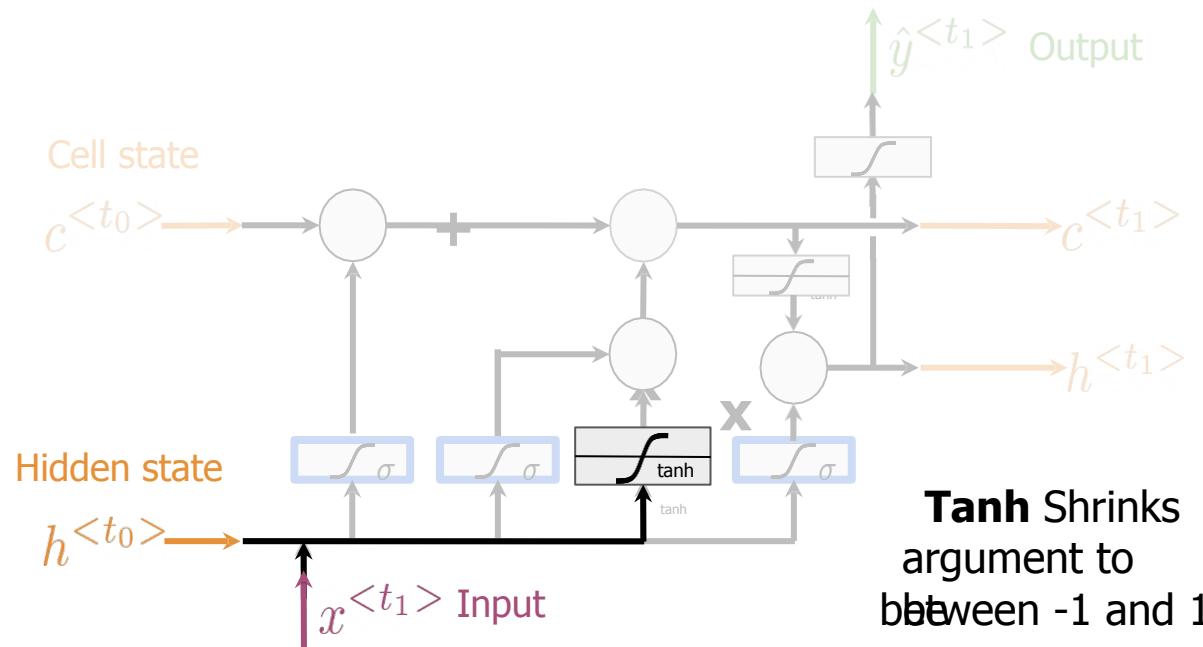
**1. Forget Gate:**  
information that is no longer important

**2. Input Gate:**  
information to be stored

**3. Output Gate:**  
information to use at current step



# Candidate Cell State



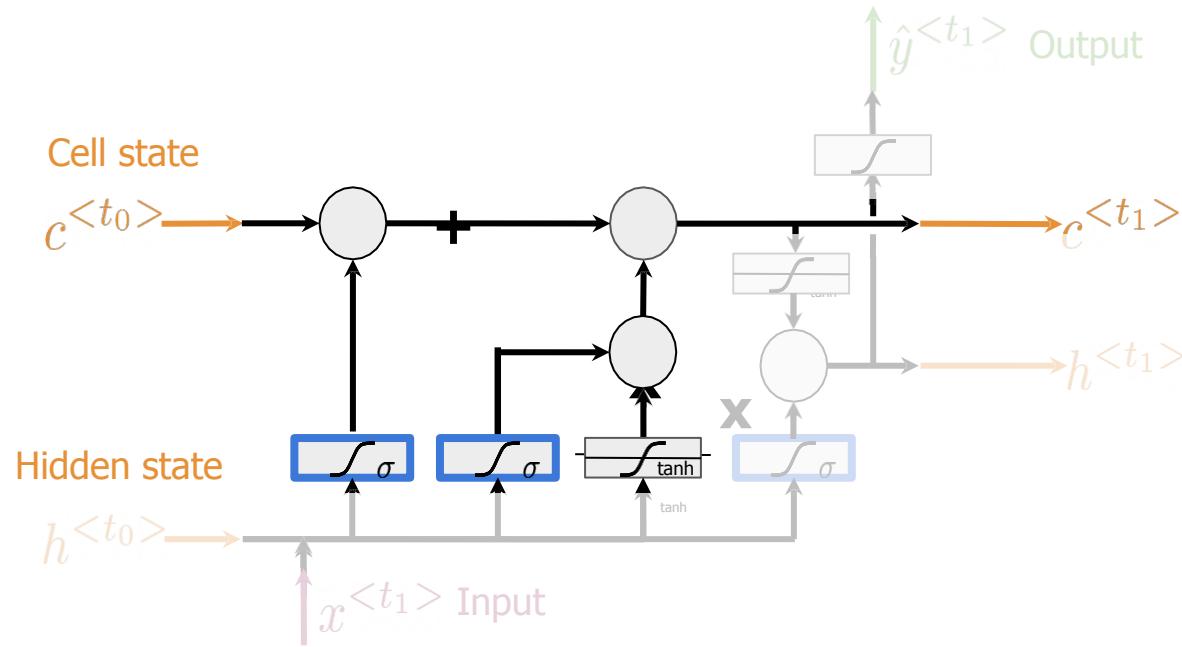
**Tanh** Shrinks argument to between -1 and 1

## Candidate cell state

Information from the previous **hidden state** and current **input**

→ **Other activations could be used**

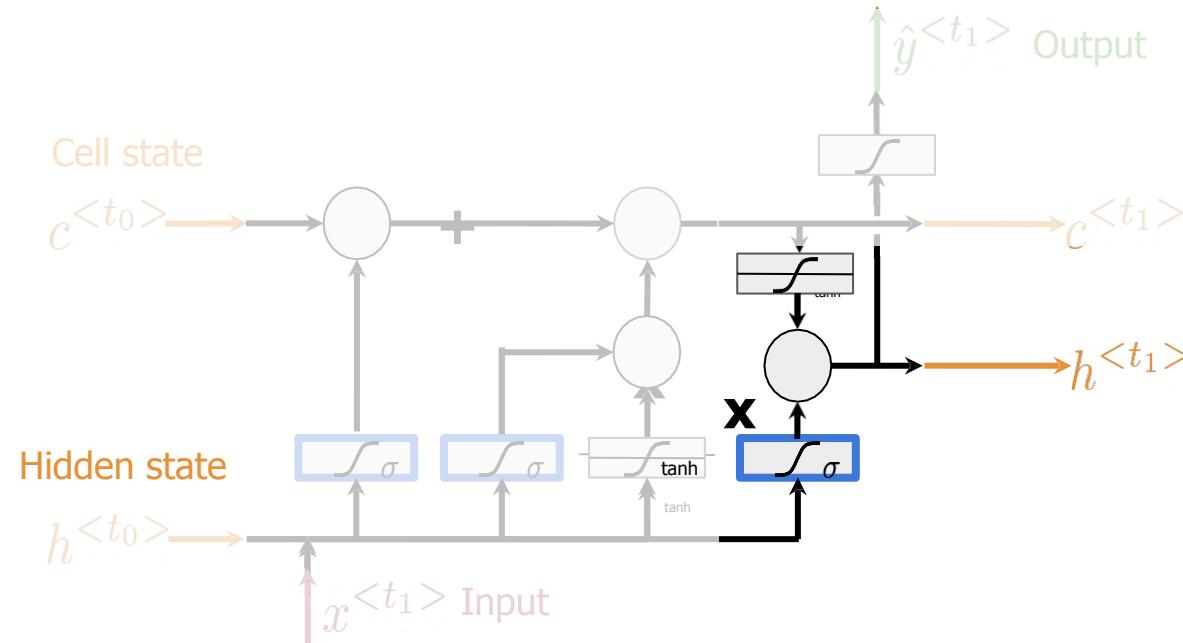
# New Cell State



## New Cell state

Add information from the **candidate cell state** using the **forget** and **input gates**

# New Hidden State



## New Hidden State

Select information from the **new cell state** using the **output gate**

The **Tanh** activation could be omitted

# Applications of LSTMs

Next-character  
prediction



Chatbots



Music  
composition



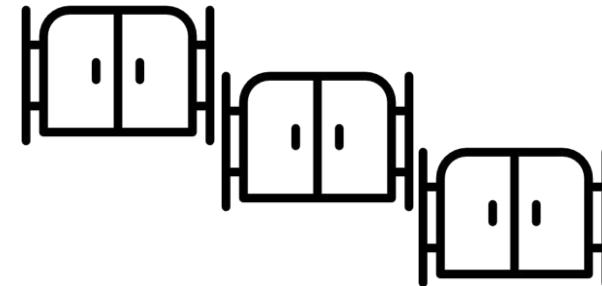
Image  
captioning



Speech  
recognition

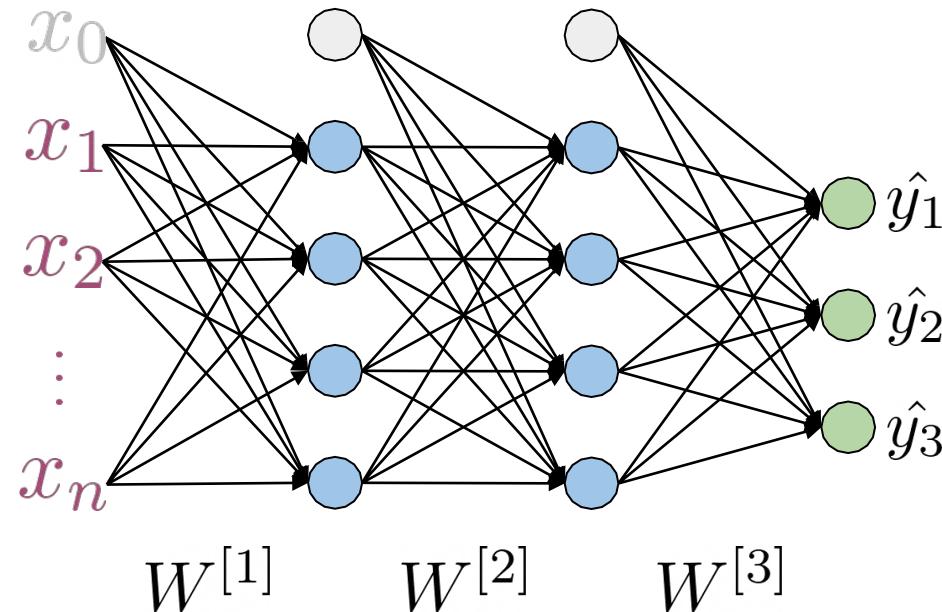


- LSTMs offer a solution to vanishing gradients
- Typical LSTMs have a cell and three gates:
  - Forget gate
  - Input gate
  - Output gate



- LSTMs use a series of gates to decide which information to keep:
  - Forget gate decides what to keep
  - Input gate decides what to add
  - Output gate decides what the next hidden state will be

# Cross Entropy Loss



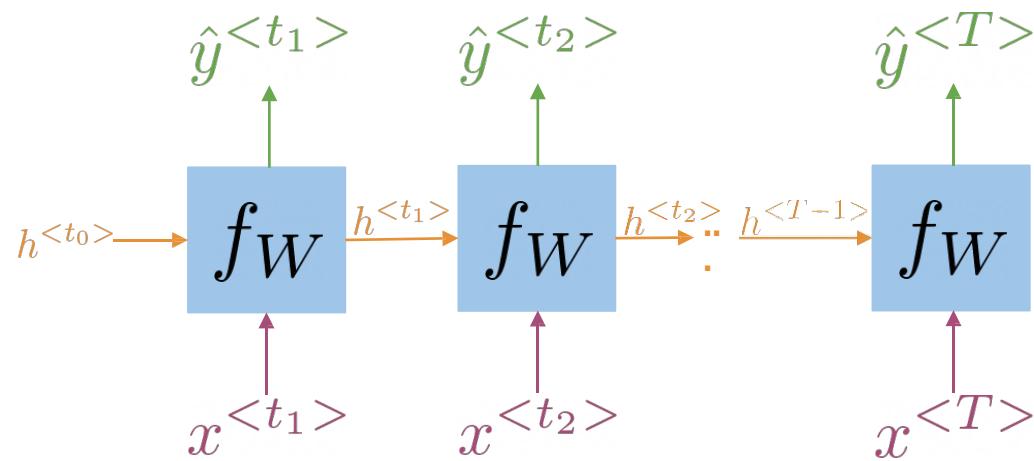
K - classes or possibilities

$$J = - \sum_{j=1}^K y_j \log \hat{y}_j$$

Either 0 or 1

Looking at a single example  $(x, y)$

# Cross Entropy Loss for RNNs



$$h^{<t>} = g(W_h[h^{<t-1>}, x^{<t>}] + b_h)$$

$$\hat{y}^{<t>} = g(W_{yh}h^{<t>} + b_y)$$

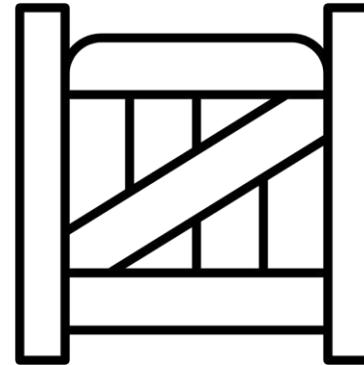
$$J = -\frac{1}{T} \sum_{t=1}^T \sum_{j=1}^K y_j^{<t>} \log \hat{y}_j^{<t>}$$

Average with respect to time

For RNNs the loss function is just an average through time!

# Gated Recurrent Unit (GRU)

- Gated recurrent unit (GRU) structure
- Comparison between GRUs and vanilla RNNs



Introduced by Cho et al., 2014

**Simpler than LSTM**, with fewer gates

- ▶ No separate memory cell
- ▶ Combines forget and input into **update gate**

**Key Components:**

- ▶ **Update Gate**  $z_t$
- ▶ **Reset Gate**  $r_t$

## Equations:

$$z_t = \sigma(W_z[x_t, h_{t-1}])$$

$$r_t = \sigma(W_r[x_t, h_{t-1}])$$

$$\tilde{h}_t = \tanh(W_h[x_t, r_t * h_{t-1}])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Comparable performance to LSTM  
Faster training, fewer parameters

# Gated Recurrent Units

“Ants are really interesting. They are everywhere.”

↓  
Plural

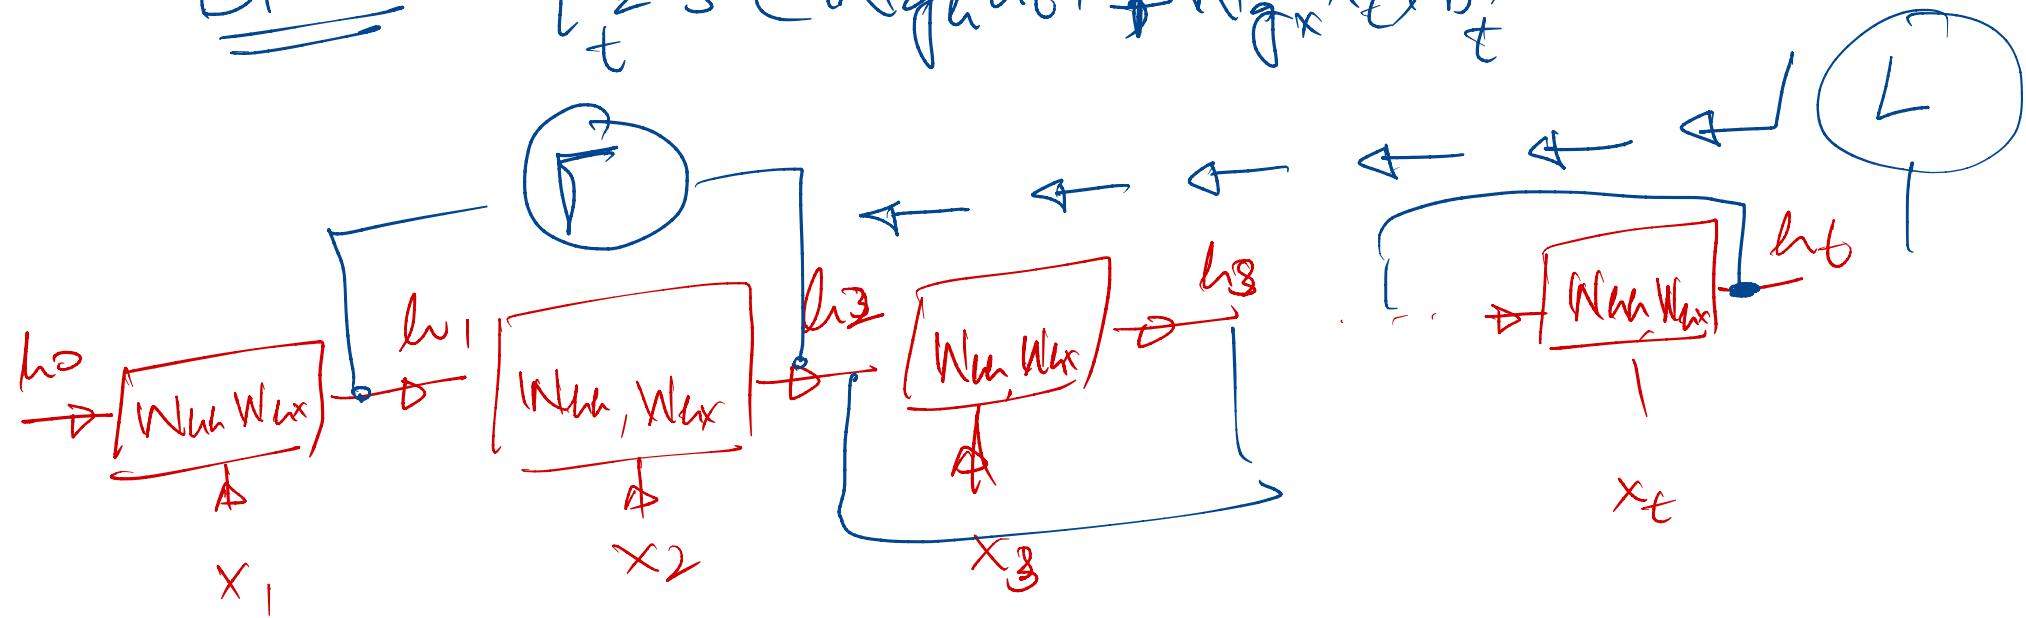
Relevance and update gates to remember important  
prior information

Vanilla

$$h_t = \tanh(W_{uh} h_{t-1} + W_{ux} x_t + b_u)$$

BPTT

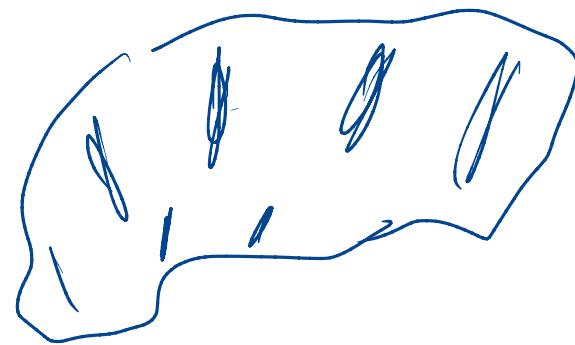
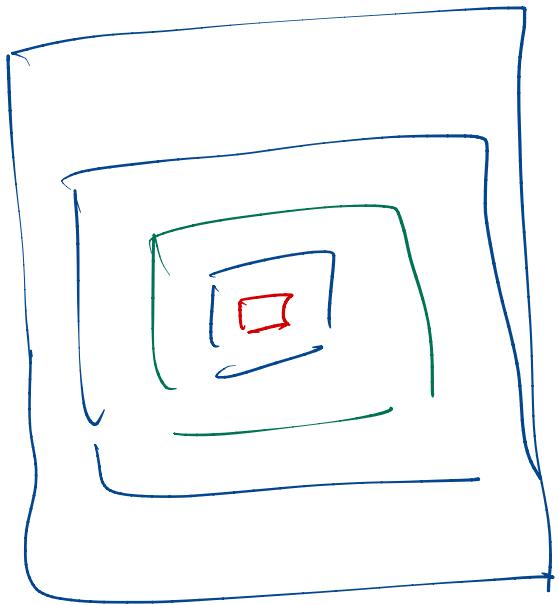
$$\sum_t^2 s (W_{gh} h_t + W_{gx} x_t + b_t)$$



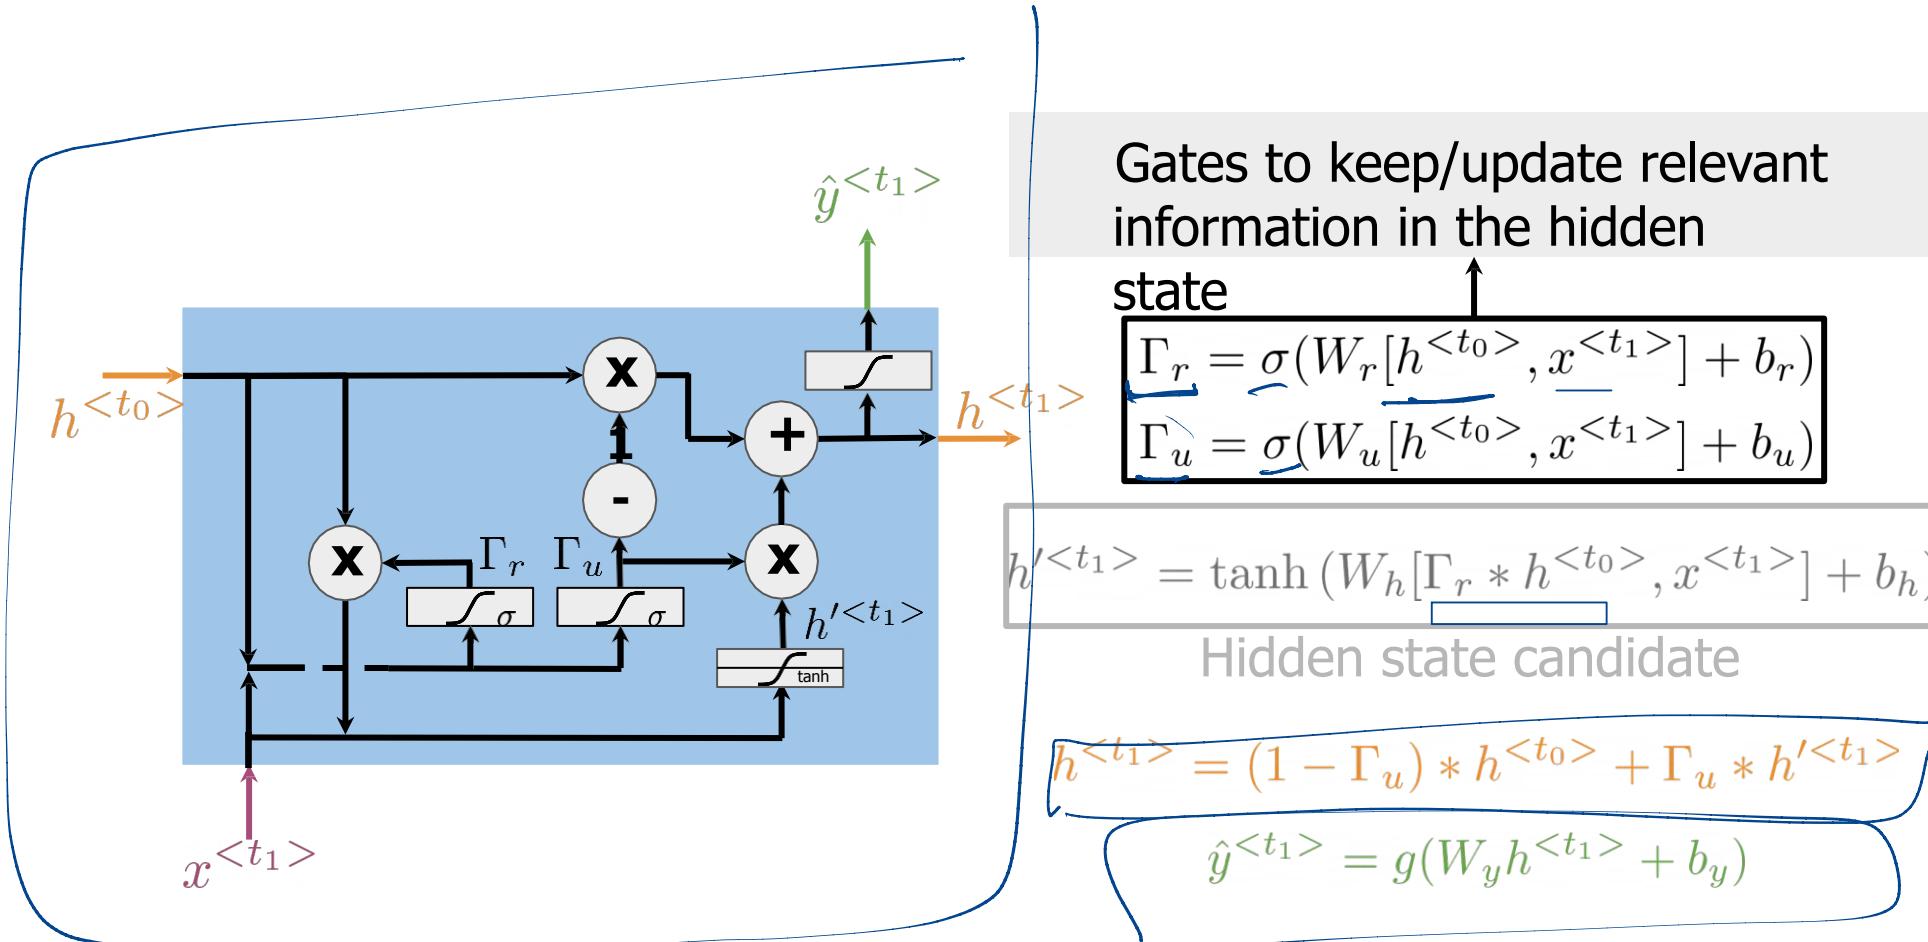
$(W_{\text{kin}})$ <sup>n</sup>  $\xrightarrow{C}$

$n \gg 1$

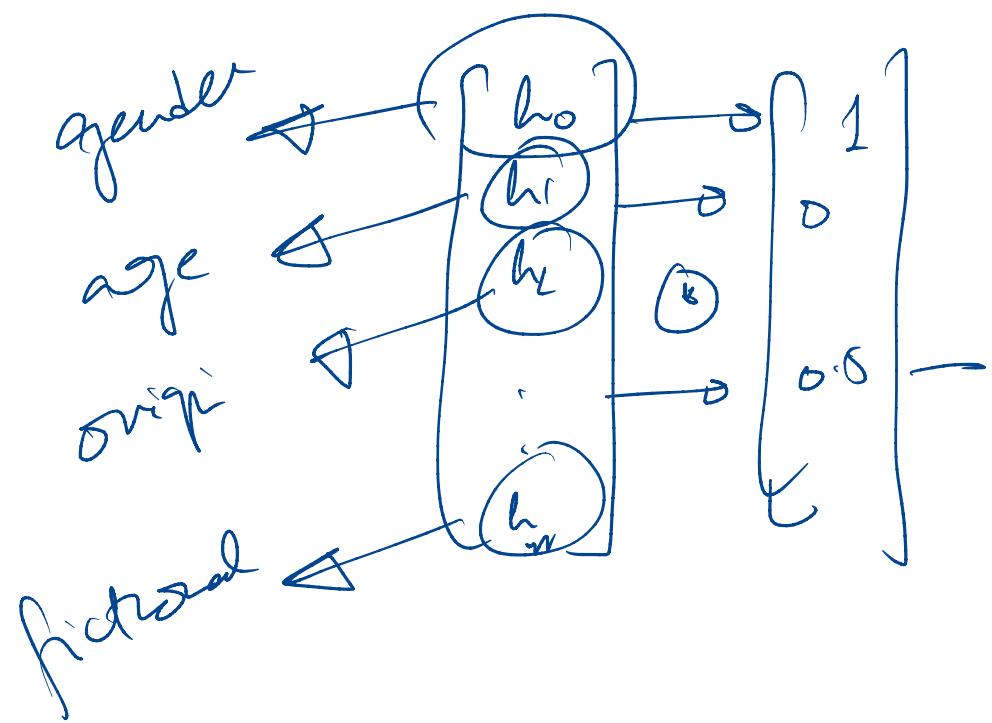
skip connections ??



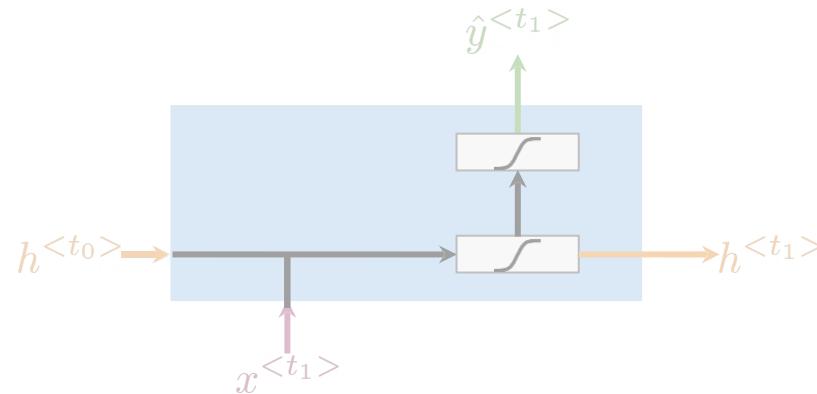
# Gated Recurrent Unit



$$\tau_w = \boxed{a(6) - 1}$$

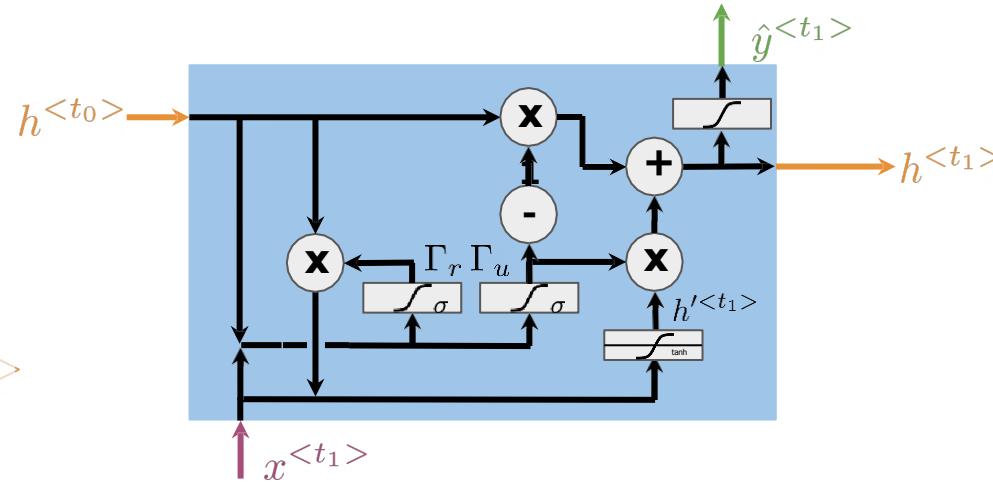


# Vanilla RNN vs GRUs



$$h^{t_1} = g(W_h[h^{t_0}, x^{t_1}] + b_h)$$

$$\hat{y}^{t_1} = g(W_y h^{t_1} + b_y)$$



$$\Gamma_u = \sigma(W_u[h^{t_0}, x^{t_1}] + b_u)$$

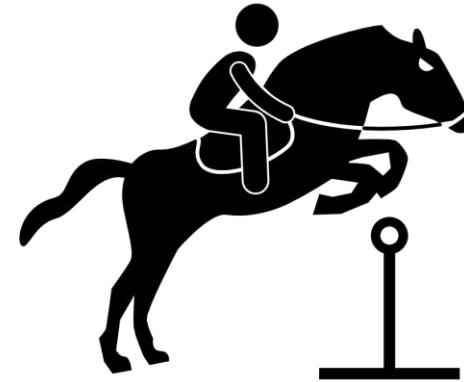
$$\Gamma_r = \sigma(W_r[h^{t_0}, x^{t_1}] + b_r)$$

$$h'^{t_1} = \tanh(W_h[\Gamma_r * h^{t_0}, x^{t_1}] + b_h)$$

$$h^{t_1} = (1 - \Gamma_u) * h^{t_0} + \Gamma_u * h'^{t_1}$$

$$\hat{y}^{t_1} = g(W_y h^{t_1} + b_y)$$

- GRUs “decide” how to update the hidden state
- GRUs help preserve important information



Feature	LSTM	GRU
Gates	3 (i, f, o)	2 (z, r)
Memory Cell	Yes	No
Complexity	Higher	Lower
Training Speed	Slower	Faster
Performance	Great for long sequences	Similar or better on short tasks

**Tip:** Try GRU first for faster results; switch to LSTM if performance suffers.

## NLP Tasks

- ▶ Language Modeling
- ▶ Machine Translation
- ▶ Sentiment Analysis
- ▶ Chatbots

## Audio & Time Series

- ▶ Music Generation
- ▶ Speech Recognition
- ▶ Anomaly Detection

## Video & Sequential Vision

- ▶ Action Recognition
- ▶ Video Captioning

Even with GRU/LSTM:

- ▶ Still sequential → Hard to parallelize
- ▶ Struggle with very long-range dependencies
- ▶ Architectural complexity
- ▶ Hard to interpret gate decisions
- ▶ Require lots of training data

- ▶ **Transformers:** Fully parallelized sequence modeling using attention
- ▶ **Efficient Attention:** Longformer, Linformer, etc. for long sequences
- ▶ **Neural Memory Networks:** Explicit memory read/write
- ▶ **Recurrent Attention Models**
- ▶ **Hybrid Architectures:** RNN + CNN + Attention

RNNs are still used in edge devices for efficient modeling

- ▶ RNNs struggle with long dependencies due to vanishing/exploding gradients
- ▶ GRU and LSTM improve memory retention using gating mechanisms
- ▶ GRU is simpler and faster; LSTM is more expressive
- ▶ Attention and transformers now dominate, but RNNs remain relevant in many domains

## Foundational Papers:

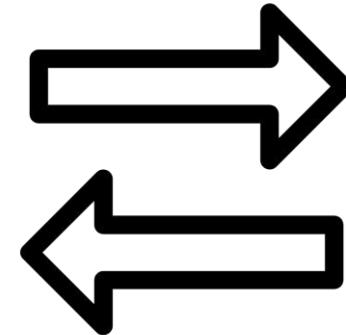
- ▶ Hochreiter, S., & Schmidhuber, J. (1997). *Long Short-Term Memory*. *Neural Computation*.
- ▶ Cho, K., van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., & Bengio, Y. (2014). *Learning Phrase Representations using RNN Encoder–Decoder with GRU*. EMNLP.
- ▶ Pascanu, R., Mikolov, T., & Bengio, Y. (2013). *On the difficulty of training RNNs*. ICML.
- ▶ Bengio, Y., Simard, P., & Frasconi, P. (1994). *Learning long-term dependencies*. IEEE Transactions on Neural Networks.
- ▶ Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). *Attention Is All You Need*. NeurIPS.

## Resources:

- ▶ Karpathy's RNN Blog:  
<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- ▶ CS231n Lecture Notes on RNNs and LSTM
- ▶ DeepLearning.ai NLP Specialization – Coursera
- ▶ MIT 6.S191 Deep Learning Lecture Slides

# Deep and Bi- directional RNNs

- How bidirectional RNNs propagate information
- Forward propagation in deep RNNs

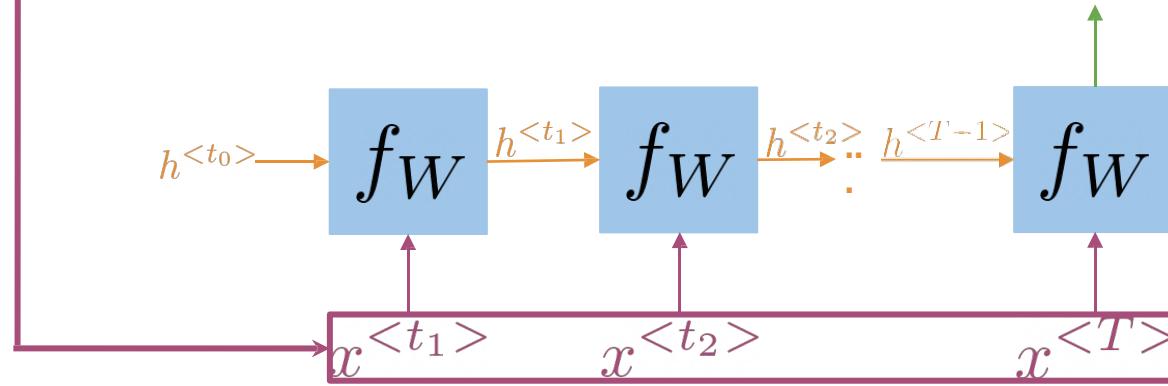


# Bi-directional RNNs

I was trying really hard to get a hold  
of his when I was about to give  
up.

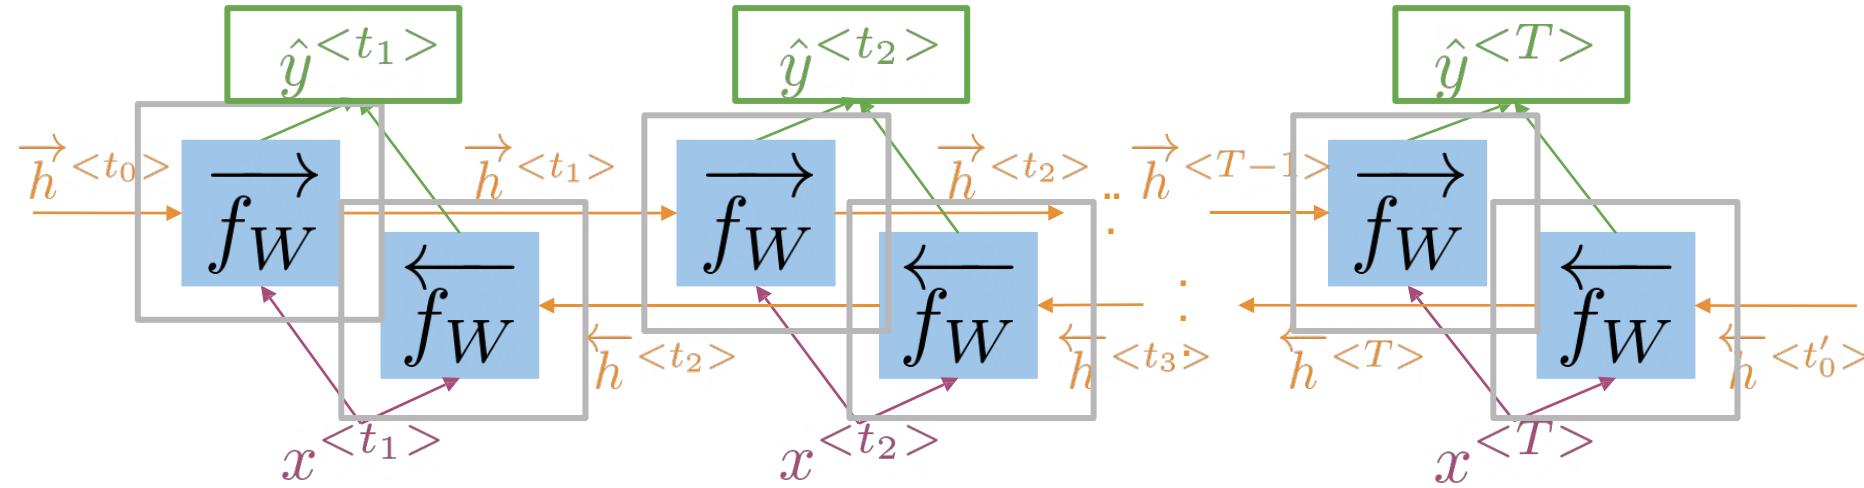
. Louise, finally

her him them



deeplearning.ai

# Bi-directional RNNs

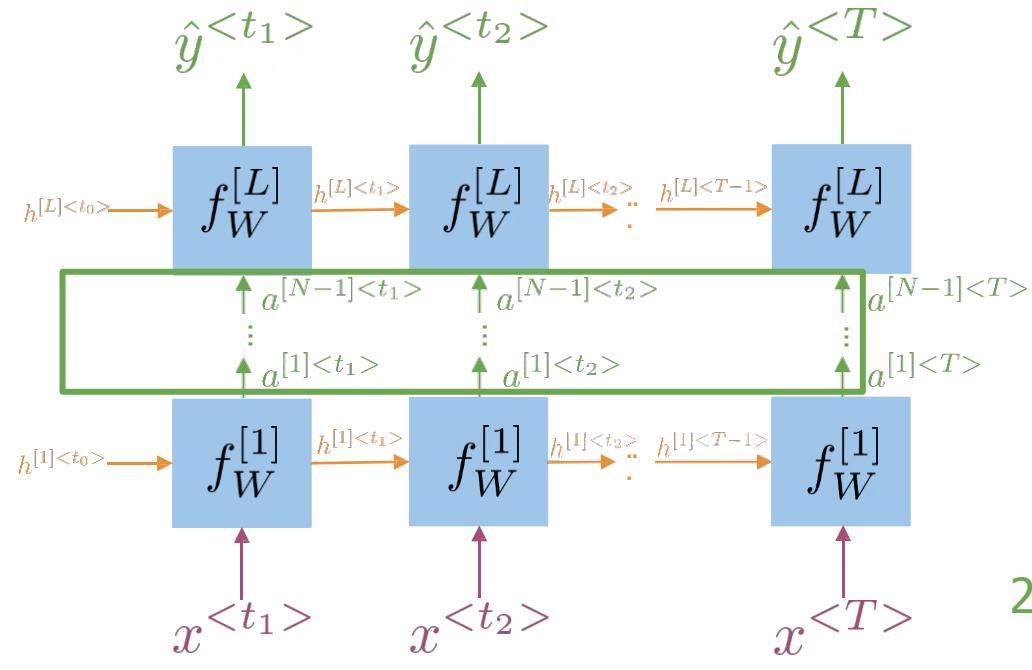


Information flows from the past and from the  
future

**independently**

$$\hat{y}^{<t>} = g(W_y[\vec{h}^{<t>}, \underline{h}^{<t>}] + b_y)$$

# Deep RNNs



$$h^{[l]} < t > = f^{[l]}(W_h^{[l]} h^{[l]} < t-1 >, a^{[l-1]} < t > + b_h^{[l]})$$

$$a^{[l]} < t > = f^{[l]}(W_a^{[l]} h^{[l]} < t > + b_a^{[l]})$$

Intermediat  
e layers  
and

1. Get hidden states for current layer
2. Pass the activations to the next layer

- In bidirectional RNNs, the outputs take information from the past and the future
- Deep RNNs have more than one layer, which helps in complex tasks



# References

These slides have been adapted from

- Younes Mourri & Lukasz Kaiser, [Natural Language Processing Specialization, DeepLearning.Ai](#)